The Journal of Open Source Software

metrica: an R package to evaluate prediction performance of regression and classification point-forecast models

Adrian A. Correndo [1¶], **Luiz H. Moro Rosso** [2], **Carlos H. Hernandez** [1], **Leonardo M. Bastos** [3], **Luciana Nieto** [1], **Dean Holzworth**[4], **and Ignacio A. Ciampitti** [1]

1 Department of Agronomy, Kansas State University, Manhattan, KS, USA. **2** Private Consultant, Brasil. **3** Department of Crop and Soil Sciences, University of Georgia, Athens, GA, USA. **4** CSIRO Agriculture and Food, Australia. ¶ Corresponding author

DOI: 10.xxxxxx/draft

Software

- Review
- Repository
- Archive

Editor: Open Journals

Reviewers:

- @openjournals

Submitted: 01 January 1970
Published: unpublished

Summary

The *metrica* R package (Correndo et al., 2022) is an open-source software designed to facilitate the quantitative and visual assessment of prediction performance of point-forecast simulation models for continuous (regression) and categorical variables (classification). The package ensembles a series of 80+ functions that account for multiple aspects of the agreement between predicted and observed values. Without the need of advanced skills on programming, *metrica* enables users to automate the estimation of multiple prediction performance metrics including goodness of fit, error metrics, error decomposition, model efficiency, indices of agreement, and to produce stylish data visualization outputs. This article introduces *metrica*, an R package developed with the main objective of contributing to transparent and reproducible evaluation of point-forecast models performance.

Statement of need

Evaluating the prediction quality is a crucial step for any simulation model, for which a myriad of metrics and visualization techniques have been developed (Tedeschi, 2006; Wallach et al., 2019; Yang et al., 2014). Nonetheless, to conduct a comprehensive assessment of the predicted-observed agreement in R (R Core Team, 2021), users normally have to rely on multiple packages, and even on self-defined functions, which increases the risk of involuntary mistakes due to the need of fluctuating syntax and data wrangling.

As the reproducibility of data analysis continues to be a challenge for science (Seibold, 2022), developing open source software like *metrica* offers a step toward a transparent and reproducible process to assist researchers in evaluating models performance. We decided to create *metrica* in R (R Core Team, 2021) due to its substantial role in data science (Thieme, 2018). Under its open-source philosophy, R empowers the democratization of statistical

computing (Hackenberger, 2020) by hosting and globally distributing cutting-edge algorithms through the Comprehensive R Archive Network (CRAN).

Finally, it is noteworthy that in the area of agricultural sciences, although point-forecast simulation models such as the Agricultural Production Systems sIMulator (APSIM) (D. Holzworth et al., 2018; D. P. Holzworth et al., 2014) count with tools to facilitate the integration into R through packages such as apsimx (Miguez, 2022), the assessment of its prediction quality is not yet integrated for R users. Therefore, we aim for *metrica* to offer users of simulation models for agriculture, plant, and soil sciences community a toolbox for assessing the performance of regression and classification point-forecast models.

Package features

For regression models, `metrica` includes four plotting functions (scatter, tiles, density, & Bland-Altman plots) using `ggplot2` (Wickham, 2016), and 48 prediction performance metrics. For classification models (two-class or multi-class), it includes one function to visualize a confusion matrix, and 27 functions of prediction scores. The full list of metrics with description, formula, and literature sources is presented in the package documentation at:

- Regression metrics: https://adriancorrendo.github.io/metrica/articles/available_metrics_regression.html.
- Classification metrics: https://adriancorrendo.github.io/metrica/articles/available_metrics_classification.html.

To extent of our knowledge, *metrica* covers several functions not supported, or partially supported by similar R packages (or components) designed for model evaluation such as `yardstick` (Kuhn & Vaughan, 2022) from `tidymodels` (Kuhn & Wickham, 2020), the measuring performance components from `caret` (Kuhn, 2022) or `mlr3` (Lang et al., 2019), `Metrics` (Hamner & Frasco, 2018), `hydroGOF` (Zambrano-Bigiarini, 2020), `cvms` (Olsen & Zachariae, 2021), `scoringutils`(Bosse et al., 2020), or `performance` (Lüdecke et al., 2021). Unique features include:

- the most extensive collection of prediction performance metrics for regression and classification models up to date.
- working under both vectorized (calling variables with $) or tabulated forms (Wickham et al., 2019).
- controlling the output format as a list (`tidy = FALSE`) or as a table (`tidy = TRUE`).
- for classification, functions automatically recognizing two-class or multi-class data; and specifically for multi-class cases, several metrics can be estimated for each class (`atom = TRUE`)(Ferri et al., 2009), (Ben-David, 2007), including balanced and imbalanced scenarios (Kubat et al., 1997).
- for regression, implementing a symmetric linear regression (standardized major axis-SMA-, (Warton et al., 2006)) to describe: i) pattern of the bivariate relationship with linear parameters (`B0_sma`, `B1_sma`), and ii) degree of predicted-observed agreement by using SMA-line to decompose the mean-squared-error (MSE) into lack of accuracy (`MLA`, `PLA`, `RMLA`) and lack of precision (`MLP`, `PLP`, `RMLP`) components (Correndo et al., 2021).
- offering MSE decomposition approaches described by (Kobayashi & Salam, 2000) (`SB`, `SDSD`, `LCS`), and (Smith & Rose, 1995) (`Ub`, `Uc`, `Ue`).
- including multiple indices of agreement and model efficiency such as: i) index of agreement `d` (Willmott, 1981), and its modified `d1` (Willmott et al., 1985) and refined `d1r` (Willmott et al., 2012) variants, ii) Nash–Sutcliffe model efficiency (`NSE`) (Nash & Sutcliffe, 1970) and its improved variants `E1` (Legates & McCabe Jr., 1999), `Erel` (Krause et al., 2005),

and Kling-Gupta model efficiency (KGE) (Kling et al., 2012), iii) Robinson's index of agreement (RAC) (Robinson, 1957, 1959), iv) Ji & Gallo agreement coefficient (AC) (Ji & Gallo, 2006), v) Duvellier's lambda (Duveiller & Meroni, 2016), vi) distance correlation (dcorr) (Székely et al., 2007), or vii) maximal information coefficient (MIC) (Reshef et al., 2011)), among others.

- importing files from APSIM Classic with `import_apsim_out()`), and from APSIM Next Generation with the `import_apsim_db()` function.

System requirements and installation

Since *metrica* operates within R, the first step is to install R (>= 4.2.0). To install and load the package:

```
# Stable version (CRAN)
install.packages("metrica")

# Development version (GitHub)
devtools::install_github("adriancorrendo/metrica")

# Load
library(metrica)
```

Using the functions

There are two core arguments to all *metrica* functions: (i) obs(Oi; observed, a.k.a. actual, measured, truth, target, label), and (ii) pred (Pi; predicted, a.k.a. simulated, fitted, modeled, estimate) values. For regression, specific functions require defining the axis `orientation` (*e.g.* predicted vs. observed -PO- or observed vs. predicted -OP-).

For two-class models, the `pos_level` argument serves to indicate the alphanumeric order of the "positive level". Following most two-class denominations as c(0,1), c("Negative", "Positive"), and c("FALSE", "TRUE"), the default `pos_level = 2` (1, "Positive", "TRUE"). However, we recognize other cases as possible (e.g. c("Crop", "NoCrop")), for which the user needs to specify `pos_level = 1`. For multi-class classification, some functions present the `atom` argument (TRUE / FALSE), which controls the output to be an overall average estimate across all classes (default), or class-wise.

Example 1: Regression (continuous variables)

The following lines of code serve to run basic regression performance analysis using a native dataset called `wheat`.

```
# Define dataset
data_wheat <- metrica::wheat

# Estimate Root Mean Square Error, result as a list
RMSE(data = data_wheat, obs = obs, pred = pred, tidy = FALSE)
#> $RMSE
#> [1] 1.666441

# Store results as a data frame
RMSE(data = data_wheat, obs = obs, pred = pred, tidy = TRUE)

#>      Metric          Score
#> 1      RMSE     1.66644142
```

To estimate multiple regression metrics at once using `metrica::metrics_summary()`:

```r
# Define metrics list
my_reg_metrics <- c("R2", "CCC", "MBE", "RMSE", "RSR", "NSE", "KGE")

# Run metrics summary
metrics_summary(data = data_wheat,
                obs = obs, pred = pred,
                type = "regression",
                metrics_list = my_reg_metrics)

#>      Metric          Score
#> 1        R2     0.84555376
#> 2       CCC     0.91553253
#> 3       MBE     0.31815953
#> 4      RMSE     1.66644142
#> 5     RRMSE     0.19094834
#> 6       RSR     0.09678632
#> 7       PLP     5.15949064
#> 8       PLA    94.84050936
#> 9       NSE     0.83871126
#> 10      KGE     0.91064709
```

To produce a scatter plot of predicted vs. observed values as a customizable `ggplot` object:

```r
scatter_plot(data = data_wheat,
             obs = obs , pred = pred,
             orientation = "PO",
             print_metrics = TRUE,
             metrics_list = my_reg_metrics,
             print_eq = TRUE,
             position_eq = c(x=14, y = 2),
             # Optional arguments to customize the plot
             shape_type = 21,
             shape_color = "steelblue",
             shape_size = 3,
             regline_type = "F1",
             regline_color = "#9e0059",
             regline_size = 2)+
  # Customize axis breaks
  scale_y_continuous(breaks = seq(0,20, by = 2))+
  scale_x_continuous(breaks = seq(0,20, by = 2))
```

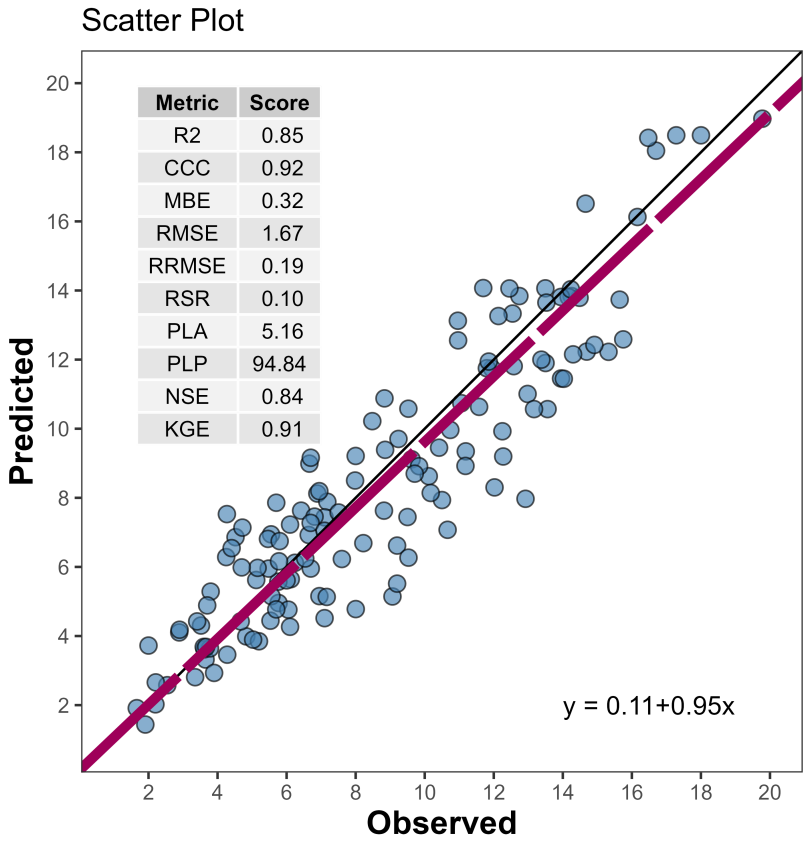

Figure 1: Predicted vs. Observed scatter plot using metrica::scatter_plot().

Example 2: Classification (categorical variables)

The following lines of code serve to run a basic classification performance analysis using a native dataset called `maize_phenology`.

```
# Define dataset
data_multiclass <- metrica::maize_phenology

# Estimate accuracy, result as a list
accuracy(data = data_multiclass, obs = actual, pred = predicted, tidy = FALSE)

#> $accuracy
#> [1] 0.8834951

# Result as a data frame
accuracy(data = data_multiclass, obs = actual, pred = predicted, tidy = TRUE)

#>         Metric       Score
#> 1     accuracy   0.8834951

# Define selected metrics
my_class_metrics <- c("accuracy", "precision", "recall", "specificity",
                      "fscore", "gmean", "khat")

# Run the summary for selected metrics
metrics_summary(data = data_multiclass,
```

```
                    obs = actual, pred = predicted,
                    type = "classification")

#>         Metric        Score
#> 1     accuracy 8.834951e-01
#> 2    precision 8.335108e-01
#> 3       recall 8.405168e-01
#> 4  specificity 9.915764e-01
#> 5       fscore 8.369991e-01
#> 6          agf 8.370017e-01
#> 7        gmean 9.129275e-01
#> 8         khat 8.624527e-01
```

To produce a confusion matrix plot users may use:

```
confusion_matrix(data = data_multiclass,
                 obs = actual, pred = predicted,
                 plot = TRUE,
                 colors = c(low="grey85" , high="steelblue"),
                 unit = "count",
                 # Print metrics_summary
                 print_metrics = TRUE,
                 # List of performance metrics
                 metrics_list = my_class_metrics,
                 # Position (bottom or top)
                 position_metrics = "bottom")
```

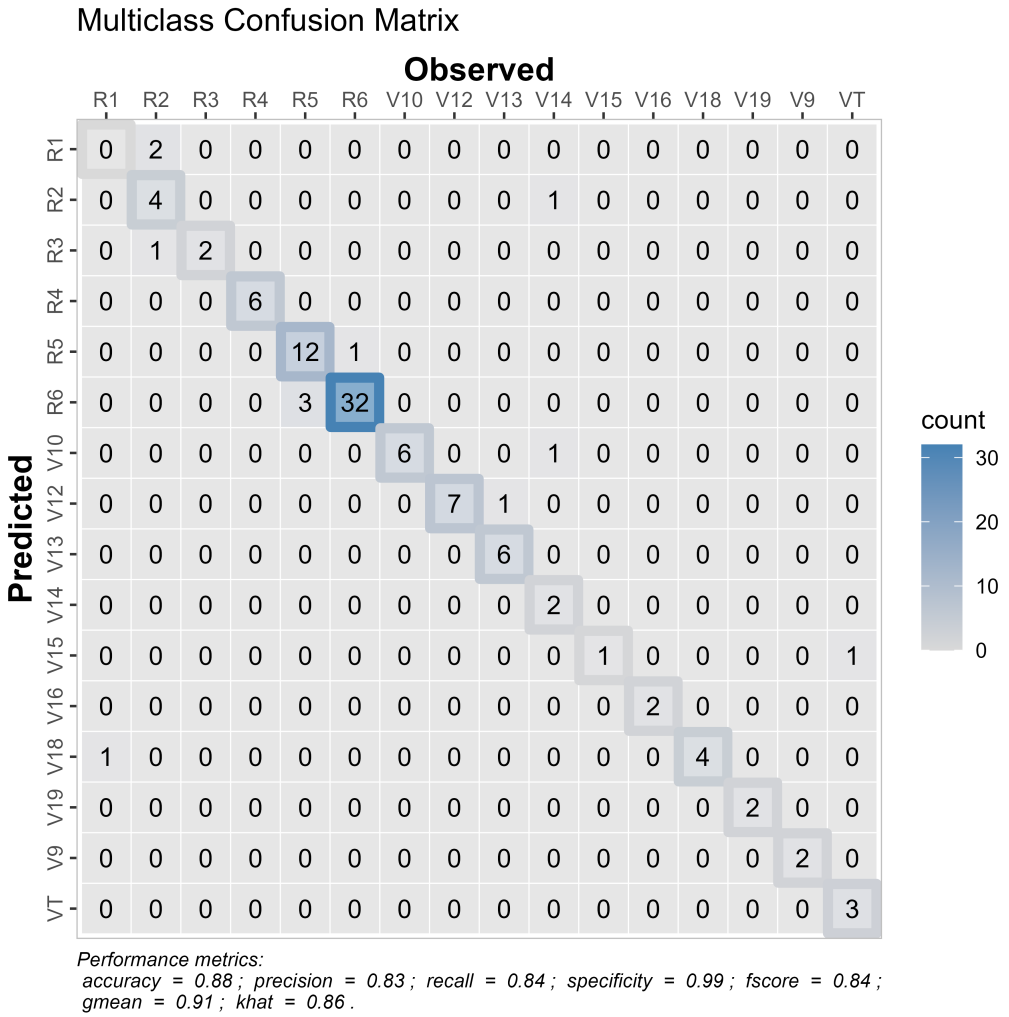

Figure 2: Confusion matrix plot using metrica::confusion_matrix().

Documentation & License

The complete documentation and vignettes of the package can be found online at https://adriancorrendo.github.io/metrica/. *metrica* is under the MIT License (https://opensource.org/licenses/MIT). Source code is available at GitHub (https://github.com/adriancorrendo/metrica) along with its corresponding section to report issues and suggestions (https://github.com/adriancorrendo/metrica/issues).

Acknowledgements

Authors gratefully acknowledge the financial support from the Feed the Future Innovation Lab for Collaborative Research on Sustainable Intensification (SIIL) at Kansas State University through funding United States Agency for International Development (USAID) under the Cooperative Agreement (Grant number AID-OAA-L-14-00006).

The Journal of Open Source Software

References

Ben-David, A. (2007). A lot of randomness is hiding in accuracy. *Engineering Applications of Artificial Intelligence*, *20*, 875–885. https://doi.org/10.1016/j.engappai.2007.01.001

Bosse, N. I., Gruson, H., Funk, S., EpiForecasts, & Abbott, S. (2020). *Scoringutils: Utilities for scoring and assessing predictions*. https://doi.org/10.5281/zenodo.4618017

Correndo, A. A., Hefley, T. J., Holzworth, D. P., & Ciampitti, I. A. (2021). Revisiting linear regression to test agreement in continuous predicted-observed datasets. *Agricultural Systems*, *192*, 103194. https://doi.org/10.1016/j.agsy.2021.103194

Correndo, A. A., Moro Rosso, L. H., Schwalbert, R., Hernandez, C., Bastos, L. M., Nieto, L., Holzworth, D., & Ciampitti, I. A. (2022). *Metrica: Prediction performance metrics*. https://CRAN.R-project.org/package=metrica

Duveiller, F., G., & Meroni, M. (2016). Revisiting the concept of a symmetric index of agreement for continuous datasets. *Scientific Reports*, *6*, 19401. https://doi.org/10.1038/srep19401

Ferri, C., Hernández-Orallo, J., & Modroiu, R. (2009). An experimental comparison of performance measures for classification. *Pattern Recognition Letters*, *30*, 27–38. https://doi.org/10.1016/j.patrec.2008.08.010

Hackenberger, B. K. (2020). R software: Unfriendly but probably the best. *Croat Med. J.*, *29;61*(1), 66–68. https://doi.org/10.3325/cmj.2020.61.66

Hamner, B., & Frasco, M. (2018). *Metrics: Evaluation metrics for machine learning*. https://CRAN.R-project.org/package=Metrics

Holzworth, D. P., Huth, N. I., deVoil, P. G., Zurcher, E. J., Herrmann, N. I., McLean, G., Chenu, K., van Oosterom, E. J., Snow, V., Murphy, C., Moore, A. D., Brown, H., Whish, J. P. M., Verrall, S., Fainges, J., Bell, L. W., Peake, A. S., Poulton, P. L., Hochman, Z., … Keating, B. A. (2014). APSIM – evolution towards a new generation of agricultural systems simulation. *Environmental Modelling & Software*, *62*, 327–350. https://doi.org/10.1016/j.envsoft.2014.07.009

Holzworth, D., Huth, N. I., Fainges, J., Brown, H., Zurcher, E., Cichota, R., Verrall, S., Herrmann, N. I., Zheng, B., & Snow, V. (2018). APSIM next generation: Overcoming challenges in modernising a farming systems model. *Environmental Modelling & Software*, *103*, 43–51. https://doi.org/10.1016/j.envsoft.2018.02.002

Ji, L., & Gallo, K. (2006). An agreement coefficient for image comparison. *Photogrammetric Engineering & Remote Sensing*, *72*(7), 823–833. https://doi.org/doi:10.14358/PERS.72.7.823

Kling, H., Fuchs, M., & Paulin, M. (2012). Runoff conditions in the upper danube basin under an ensemble of climate change scenarios. *Journal of Hydrology*, *424-425*, 264–277. https://doi.org/10.1016/j.jhydrol.2012.01.011

Kobayashi, K., & Salam, M. U. (2000). Comparing simulated and measured values using mean squared deviation and its components. *Agronomy Journal*, *92*(2), 345–352. https://doi.org/10.2134/agronj2000.922345x

Krause, P., Boyle, D. P., & Bäse, F. (2005). Comparison of different efficiency criteria for hydrological model assessment. *Advances in Geosciences*, *5*, 89–97. https://doi.org/10.5194/adgeo-5-89-2005

Kubat, M., Matwin, S., & others. (1997). Addressing the curse of imbalanced training sets: One-sided selection. *Icml*, *97*, 179.

Kuhn, M. (2022). *Caret: Classification and regression training*. https://CRAN.R-project.org/package=caret

Kuhn, M., & Vaughan, D. (2022). *Yardstick: Tidy characterizations of model performance*. https://CRAN.R-project.org/package=yardstick

Kuhn, M., & Wickham, H. (2020). *Tidymodels: A collection of packages for modeling and machine learning using tidyverse principles*. https://www.tidymodels.org

Lang, M., Binder, M., Richter, J., Schratz, P., Pfisterer, F., Coors, S., Au, Q., Casalicchio, G., Kotthoff, L., & Bischl, B. (2019). mlr3: A modern object-oriented machine learning framework in R. *Journal of Open Source Software*. https://doi.org/10.21105/joss.01903

Legates, D. R., & McCabe Jr., G. J. (1999). Evaluating the use of "goodness-of-fit" measures in hydrologic and hydroclimatic model validation. *Water Resources Research*, *35*(1), 233–241. https://doi.org/10.1029/1998WR900018

Lüdecke, D., Ben-Shachar, M. S., Patil, I., Waggoner, P., & Makowski, D. (2021). performance: An R package for assessment, comparison and testing of statistical models. *Journal of Open Source Software*, *6*(60), 3139. https://doi.org/10.21105/joss.03139

Miguez, F. (2022). *Apsimx: Inspect, read, edit and run 'APSIM' "next generation" and 'APSIM' classic*. https://CRAN.R-project.org/package=apsimx

Nash, J. E., & Sutcliffe, J. V. (1970). River flow forecasting through conceptual models part i — a discussion of principles. *Journal of Hydrology*, *10*(3), 282–290. https://doi.org/10.1016/0022-1694(70)90255-6

Olsen, L. R., & Zachariae, H. B. (2021). *Cvms: Cross-validation for model selection*. https://CRAN.R-project.org/package=cvms

R Core Team. (2021). *R: A language and environment for statistical computing*. R Foundation for Statistical Computing. https://www.R-project.org/

Reshef, D. N., Reshef, Y. A., Finucane, H. K., Grossman, S. R., McVean, G., Turnbaugh, P. J., Lander, E. S., Mitzenmacher, M., & Sabeti, P. C. (2011). Detecting novel associations in large data sets. *Science*, *334*(6062), 1518–1524. https://doi.org/10.1126/science.1205438

Robinson, W. S. (1957). The statistical measurement of agreement. *American Sociological Review*, *22*(1), 17–25. https://doi.org/10.2307/2088760

Robinson, W. S. (1959). The geometric interpretation of agreement. *American Sociological Review*, *24*(3), 338–345. https://doi.org/10.2307/2089382

Seibold, S. A. D., Heidi AND Czerny. (2022). Correction: A computational reproducibility study of PLOS ONE articles featuring longitudinal data analyses. *PLOS ONE*, *17*(5), 1–1. https://doi.org/10.1371/journal.pone.0269047

Smith, E. P., & Rose, K. A. (1995). Model goodness-of-fit analysis using regression and related techniques. *Ecological Modelling*, *77*(1), 49–64. https://doi.org/10.1016/0304-3800(93)E0074-D

Székely, G. J., Rizzo, M. L., & Bakirov, N. K. (2007). Measuring and testing dependence by correlation of distances. *The Annals of Statistics*, *35*(6), 2769–2794. https://doi.org/10.1214/009053607000000505

Tedeschi, L. O. (2006). Assessment of the adequacy of mathematical models. *Agricultural Systems*, *89*(2), 225–247. https://doi.org/10.1016/j.agsy.2005.11.004

Thieme, N. (2018). R generation. *Significance*, *15*(4), 14–19. https://doi.org/10.1111/j.1740-9713.2018.01169.x

Wallach, D., Makowski, D., Jones, J. W., & Brun, F. (2019). Chapter 9 - model evaluation. In D. Wallach, D. Makowski, J. W. Jones, & F. Brun (Eds.), *Working with dynamic*

crop models (third edition) (Third Edition, pp. 311–373). Academic Press. https://doi.org/10.1016/B978-0-12-811756-9.00009-5

Warton, D. I., Wright, I. J., Falster, D. S., & Westoby, M. (2006). Bivariate line-fitting methods for allometry. *Biological Reviews*, *81*(2), 259–291. https://doi.org/10.1017/S1464793106007007

Wickham, H. (2016). *ggplot2: Elegant graphics for data analysis*. Springer-Verlag New York. ISBN: 978-3-319-24277-4

Wickham, H., Averick, M., Bryan, J., Chang, W., McGowan, L. D., François, R., Grolemun, G., Haye, A., Henr, L., Heste, J., Kuh, M., Pederse, T. L., Mille, E., Bach, S. M., Müll, K., Jeroen Oo, David Robins, Dana Paige Seid, Vitalie Spi, … Hiroaki Yutani. (2019). Welcome to the tidyverse. *Journal of Open Source Software*, *4*(43), 1686. https://doi.org/10.21105/joss.01686

Willmott, C. J. (1981). ON THE VALIDATION OF MODELS. *Physical Geography*, *2*(2), 184–194. https://doi.org/10.1080/02723646.1981.10642213

Willmott, C. J., Ackleson, S. G., Davis, R. E., Feddema, J. J., Klink, K. M., Legates, D. R., O'Donnell, J., & Rowe, C. M. (1985). Statistics for the evaluation and comparison of models. *Journal of Geophysical Research: Oceans*, *90*(C5), 8995–9005. https://doi.org/10.1029/JC090iC05p08995

Willmott, C. J., Robeson, S. M., & Matsuura, K. (2012). A refined index of model performance. *International Journal of Climatology*, *32*(13), 2088–2094. https://doi.org/10.1002/joc.2419

Yang, J. M., Yang, J. Y., Liu, S., & Hoogenboom, G. (2014). An evaluation of the statistical methods for testing the performance of crop models with observed data. *Agricultural Systems*, *127*, 81–89. https://doi.org/10.1016/j.agsy.2014.01.008

Zambrano-Bigiarini, M. (2020). *hydroGOF: Goodness-of-fit functions for comparison of simulated and observed hydrological time series*. https://doi.org/10.5281/zenodo.839854